

Why ReLU networks yield high-confidence predictions far away from the training data and how to mitigate the problem(CVPR 2019 Oral)

Juhyeong Kim

SKKU Visual Computing Lab

wngud0811@naver.com

September 22, 2020

Overview

- 1 Introduction
- 2 ReLU networks produce piecewise affine function
- 3 Why ReLU networks produce high confidence predictions far away from the training data
- 4 Adversarial Confidence Enhanced Training
- 5 Experiments

- 1 Introduction
- 2 ReLU networks produce piecewise affine function
- 3 Why ReLU networks produce high confidence predictions far away from the training data
- 4 Adversarial Confidence Enhanced Training
- 5 Experiments

Problem

- For many popular deep learning models,
- High confidence output can be made for out-of-distribution data.
- Also, often produce over-confident predictions in original tasks.


Importance of uncertainty Quantification

- Deep learning models may fail in the case of noisy data or out-of-distribution data.

Train time



rabbit: **0.8** / wolf: 0.2



rabbit: 0.3 / wolf: **0.7**

Test time



rabbit: 0.1 / wolf: **0.9**


① Useful Solutions

- Dropout as a bayesian approximation: Representing model uncertainty in deep learning.(ICML 2016)
- On calibration of modern neural networks.(ICML 2017)
- Simple and scalable predictive uncertainty estimation using deep ensembles.(NIPS 2017)
- Also, there exist classifiers which is not being confident in areas where one has never seen data.(ex. RBF networks)

- 1 Introduction
- 2 ReLU networks produce piecewise affine function
- 3 Why ReLU networks produce high confidence predictions far away from the training data
- 4 Adversarial Confidence Enhanced Training
- 5 Experiments


- ① For a function $f : U \rightarrow V$, we call it **linear** if,
 - $f(x + y) = f(x) + f(y)$
 - $\alpha f(x) = f(\alpha x)$
 - For $\forall x, y \in U$
- ② For a function $f : U \rightarrow V$, we call it **affine** if,
 - $\exists b \in V$ such that $f = \bar{f} + b$ and \bar{f} is linear function.
- ③ Linear function \subset Affine function

Polytope



- Polygon in arbitrary dimensional space.


Polytope



- Let each red line in figure is equal to $w_{\cdot,i}^T x + b_i = 0$ at first layer.
- If the red line is separation at first layer, separation at second layer is like orange line.
- Each ReLU activation can be considered as dividing input space into two parts.

Piecewise Affine

- A function f is called **piecewise affine** if there exists finite set of polytopes $\{Q_r\}_{r=1}^M$ such that $\bigcup_{r=1}^M Q_r = \mathbb{R}^d$ and f is affine function in each Q_r , $r = 1, 2, \dots, M$.



Neural Network

- Let $W^{(l)} \in \mathbb{R}^{n_l \times n_{l-1}}$ and $b^{(l)} \in \mathbb{R}^{n_l}$ for $l = 1, 2, \dots, L$ are parameters.
- Then, ReLU Network can be expressed as:
$$f^{(k)}(x) = W^{(k)} \text{ReLU}(f^{(k-1)}) + b^{(k)}, \quad k = 1, \dots, L$$

- Define a function

$$\Sigma^{(l)}(x)_{ij} = \begin{cases} 1 & \text{if } i = j \text{ and } f_i^{(l)}(x) > 0, \\ 0 & \text{else.} \end{cases}$$

- $$\begin{aligned} f^{(k)}(x) &= W^{(k)} \Sigma^{(k-1)}(x) \left(W^{(k-1)} \Sigma^{(k-2)}(x) \right. \\ &\quad \times \left(\dots \left(W^{(1)}x + b^{(1)} \right) \dots \right) + b^{(k-1)} \Big) + b^{(k)} \\ &= V^{(k)} \cdot x + a^{(k)} \end{aligned}$$

for $k = 1, \dots, L$ with $V^{(k)} \in \mathbb{R}^{n_k \times d}$ and $a^{(k)} \in \mathbb{R}^{n_k}$

Neural Network

- $$\begin{aligned} f^{(k)}(x) &= W^{(k)} \Sigma^{(k-1)}(x) \left(W^{(k-1)} \Sigma^{(k-2)}(x) \right. \\ &\quad \times \left(\dots \left(W^{(1)} x + b^{(1)} \right) \dots \right) + b^{(k-1)} \Big) + b^{(k)} \\ &= V^{(k)} \cdot x + a^{(k)} \end{aligned}$$

for $k = 1, \dots, L$ with $V^{(k)} \in \mathbb{R}^{n_k \times d}$ and $a^{(k)} \in \mathbb{R}^{n_k}$

$$\begin{aligned} V^{(k)} &= W^{(k)} \left(\prod_{l=1}^{k-1} \Sigma^{(k-l)}(x) W^{(k-l)} \right) \\ a^{(k)} &= b^{(k)} + \sum_{l=1}^{k-1} \left(\prod_{m=1}^{k-l} W^{(k+1-m)} \Sigma^{(k-m)}(x) \right) b^{(l)} \end{aligned}$$

Conclusion

- ReLU network = piecewise affine function + softmax.
- Similarly, many activation functions, including leaky ReLU can be shown to be piecewise affine.
- Fully connected, Convolution with average/max pooling, Residual layers are included.

Implications


- Entire ReLU network is just simple softmax classifier when domain is restricted to each polytope.

- 1 Introduction
- 2 ReLU networks produce piecewise affine function
- 3 Why ReLU networks produce high confidence predictions far away from the training data
- 4 Adversarial Confidence Enhanced Training
- 5 Experiments

Lemma 3.1

Let ReLU-classifier divide input space to set of linear regions $\{Q_l\}_{l=1}^R$.

Then, any point $x \in \mathbb{R}^d$ in input space, there exist $a \in \mathbb{R}$ with $a > 0$, such that $\beta x \in Q_t$ for all $\beta \geq a$, $t \in \{1, \dots, R\}$.



Theorem 3.1

If the softmax input of ReLU network is piecewise affine function,
for almost any input x and any threshold level $0 < t < 1$,
there exists some constant $\alpha > 0$,

$$\frac{e^{f_k(\alpha x)}}{\sum_{r=1}^K e^{f_r(\alpha x)}} \geq t$$

Theorem 3.1

And also,
$$\lim_{\alpha \rightarrow \infty} \frac{e^{f_k(\alpha x)}}{\sum_{r=1}^K e^{f_r(\alpha x)}} = 1$$

Limitation

- Author assumed \mathbb{R}^d as input space, but many applications assume $[0, 1]^d$ as input space.
- In these cases, theorem does not directly applied.
- But empirically, training in bounded domain shows same problems.

- 1 Introduction
- 2 ReLU networks produce piecewise affine function
- 3 Why ReLU networks produce high confidence predictions far away from the training data
- 4 Adversarial Confidence Enhanced Training**
- 5 Experiments

- Author proposed two methods:
 - Confidence Enhancing Data Augmentation(CEDA)
 - Adversarial Confidence Enhanced Training(ACET)

Confidence Enhancing Data Augmentation(CEDA)


- In image classification problem:
 - 1 Assume the type of out-of-distribution data in domain.
(For example, Uniform distribution on $[0, 1]^{c \times w \times h}$)
 - 2 Sample random noise from distribution.
 - 3 Forward noise image to model and get softmax probabilities.
 - 4 Minimize the maximum log softmax value.

$$\Rightarrow \text{Regularizer} \quad \lambda \mathbb{E} \left[\max_{l=1, \dots, K} \log \left(\frac{e^{f_l(z)}}{\sum_{k=1}^K e^{f_l(z)}} \right) \right]$$

Adversarial Confidence Enhanced Training(ACET)

- However, CEDA may need multiple sampling and forward which will yield high computational costs.
- Inspired by adversarial training, author proposed more easier and more scalable method.
- Modified regularizer $\lambda \mathbb{E} \left[\max_{\|u-z\|_p \leq \epsilon} \max_{l=1, \dots, K} \log \left(\frac{e^{f_l(z)}}{\sum_{k=1}^K e^{f_l(z)}} \right) \right]$

- 1 Introduction
- 2 ReLU networks produce piecewise affine function
- 3 Why ReLU networks produce high confidence predictions far away from the training data
- 4 Adversarial Confidence Enhanced Training
- 5 Experiments



Trained on MNIST	Plain (TE: 0.51%)			CEDA (TE: 0.74%)			ACET (TE: 0.66%)		
	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95
MNIST	0.991	–	–	0.987	–	–	0.986	–	–
FMNIST	0.654	0.972	0.121	0.373	0.994	0.027	0.239	0.998	0.003
EMNIST	0.821	0.883	0.374	0.787	0.895	0.358	0.752	0.912	0.313
grayCIFAR-10	0.492	0.996	0.003	0.105	1.000	0.000	0.101	1.000	0.000
Noise	0.463	0.998	0.000	0.100	1.000	0.000	0.100	1.000	0.000
Adv. Noise	1.000	0.031	1.000	0.102	0.998	0.002	0.162	0.992	0.042
Adv. Samples	0.999	0.358	0.992	0.987	0.549	0.953	0.854	0.692	0.782
Trained on SVHN	Plain (TE: 3.53%)			CEDA (TE: 3.50%)			ACET (TE: 3.52%)		
	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95
SVHN	0.980	–	–	0.977	–	–	0.978	–	–
CIFAR-10	0.732	0.938	0.348	0.551	0.960	0.209	0.435	0.973	0.140
CIFAR-100	0.730	0.935	0.350	0.527	0.959	0.205	0.414	0.971	0.139
LSUN CR	0.722	0.945	0.324	0.364	0.984	0.084	0.148	0.997	0.012
Imagenet-	0.725	0.939	0.340	0.574	0.955	0.232	0.368	0.977	0.113
Noise	0.720	0.943	0.325	0.100	1.000	0.000	0.100	1.000	0.000
Adv. Noise	1.000	0.004	1.000	0.946	0.062	0.940	0.101	1.000	0.000
Adv. Samples	1.000	0.004	1.000	0.995	0.009	0.994	0.369	0.778	0.279
Trained on CIFAR-10	Plain (TE: 8.87%)			CEDA (TE: 8.87%)			ACET (TE: 8.44%)		
	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95
CIFAR-10	0.949	–	–	0.946	–	–	0.948	–	–
SVHN	0.800	0.850	0.783	0.327	0.978	0.146	0.263	0.981	0.118
CIFAR-100	0.764	0.856	0.715	0.761	0.850	0.720	0.764	0.852	0.711
LSUN CR	0.738	0.872	0.667	0.735	0.864	0.680	0.745	0.858	0.677
Imagenet-	0.757	0.858	0.698	0.749	0.853	0.704	0.744	0.859	0.678
Noise	0.825	0.827	0.818	0.100	1.000	0.000	0.100	1.000	0.000
Adv. Noise	1.000	0.035	1.000	0.985	0.032	0.983	0.112	0.999	0.008
Adv. Samples	1.000	0.034	1.000	1.000	0.014	1.000	0.633	0.512	0.590
Trained on CIFAR-100	Plain (TE: 31.97%)			CEDA (TE: 32.74%)			ACET (TE: 32.24%)		
	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95	MMC	AUROC	FPR@95
CIFAR-100	0.751	–	–	0.734	–	–	0.728	–	–
SVHN	0.570	0.710	0.865	0.290	0.874	0.410	0.234	0.912	0.345
CIFAR-10	0.560	0.718	0.856	0.547	0.711	0.855	0.530	0.720	0.860
LSUN CR	0.592	0.690	0.887	0.581	0.678	0.887	0.554	0.698	0.881
Imagenet-	0.531	0.744	0.827	0.504	0.749	0.808	0.492	0.752	0.819
Noise	0.614	0.672	0.928	0.010	1.000	0.000	0.010	1.000	0.000
Adv. Noise	1.000	0.000	1.000	0.985	0.015	0.985	0.013	0.998	0.003
Adv. Samples	0.999	0.010	1.000	0.999	0.012	1.000	0.863	0.267	0.975

Summary

- ReLU networks always make high confidence predictions far away from training data.
- Temperature rescaling and reject option does not help.
- Using modified training similar to adversarial training, we can reduce high confidence problem at out-of-distribution data.